

Logistic Regression Using R

To understand Logistic Regression using R, we will use
“menarche” data and “mtcars” data

In the built-in data set `mtcars`, (MASS, library) the data column `am` represents the transmission type of the automobile model (0 = automatic, 1 = manual). With the logistic regression equation, we can model the probability of a manual transmission in a vehicle based on its engine horsepower and weight data.

```
> attach(mtcars)
> head(mtcars)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1

```
> |
```

```
> am.glm = glm(formula=am ~ hp + wt, data=mtcars, family=binomial)
>
> ###We then wrap the test parameters inside a data frame newdata.
>
> summary(am.glm)
```

Call:

```
glm(formula = am ~ hp + wt, family = binomial, data = mtcars)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.2537	-0.1568	-0.0168	0.1543	1.3449

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	18.86630	7.44356	2.535	0.01126 *
hp	0.03626	0.01773	2.044	0.04091 *
wt	-8.08348	3.06868	-2.634	0.00843 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 43.230 on 31 degrees of freedom
Residual deviance: 10.059 on 29 degrees of freedom
AIC: 16.059

Number of Fisher Scoring iterations: 8

```

>
> hp=c(100, 105, 110, 115, 120)
> wt=c(2.8, 2, 2.4, 2.2, 2.6)
> newdata = data.frame(hp, wt)
> pred=predict(am.glm, newdata, type="response")
> pred
      1      2      3      4      5
0.4645913 0.9985081 0.9693524 0.9947905 0.9002437
>
>
>
> out=cbind(hp, wt, pred)
> out
   hp  wt   pred
1 100 2.8 0.4645913
2 105 2.0 0.9985081
3 110 2.4 0.9693524
4 115 2.2 0.9947905
5 120 2.6 0.9002437
> |

```

```

> hp=mtcars$hp[1:10]
> wt=mtcars$wt[1:10]
> newdata = data.frame(hp, wt)
> pred=predict(am.glm, newdata, type="response")
> pred
      1      2      3      4      5
0.842335537 0.404782533 0.970240822 0.041728035 0.069388122
      6      7      8      9     10
0.004988159 0.248041206 0.009265579 0.040998134 0.011190709
> out=cbind(hp, wt, am, pred)
> out
   hp   wt am   pred
1 110 2.620 1 0.842335537
2 110 2.875 1 0.404782533
3  93 2.320 1 0.970240822
4 110 3.215 0 0.041728035
5 175 3.440 0 0.069388122
6 105 3.460 0 0.004988159
7 245 3.570 0 0.248041206
8  62 3.190 0 0.009265579
9  95 3.150 0 0.040998134
10 123 3.440 0 0.011190709
> |

```

```
>  
> with(am.glm, null.deviance - deviance)  
[1] 33.17062  
> with(am.glm, df.null - df.residual)  
[1] 2  
> with(am.glm, pchisq(null.deviance - deviance, df.null - df.residual, lower.tail = FALSE))  
[1] 6.267449e-08  
>  
> |
```

R Console

```
> am.glm = glm(formula=am ~ hp + wt + disp + gear, data=mtcars, fami$
Warning messages:
1: glm.fit: algorithm did not converge
2: glm.fit: fitted probabilities numerically 0 or 1 occurred
> ###We then wrap the test parameters inside a data frame newdata.
> summary(am.glm)
```

Call:

```
glm(formula = am ~ hp + wt + disp + gear, family = binomial,
     data = mtcars)
```

Deviance Residuals:

	Min	1Q	Median	3Q	Max
	-2.996e-05	-2.110e-08	-2.110e-08	2.110e-08	2.501e-05

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-4.248e+01	4.616e+05	0.000	1.000
hp	4.186e-01	2.075e+03	0.000	1.000
wt	-1.165e+02	1.479e+05	-0.001	0.999
disp	2.815e-01	1.674e+03	0.000	1.000
gear	7.717e+01	9.730e+04	0.001	0.999

(Dispersion parameter for binomial family taken to be 1)

```
Null deviance: 4.3230e+01 on 31 degrees of freedom
Residual deviance: 2.2394e-09 on 27 degrees of freedom
AIC: 10
```



```
> hp=mtcars$hp[1:10]
> wt=mtcars$wt[1:10]
> disp=mtcars$disp[1:10]
> gear=mtcars$gear[1:10]
> am=mtcars$am[1:10]
> pred=predict(am.glm, newdata, type="response")
> pred
      1      2      3      4      5      6      7
1.000000e+00 1.000000e+00 1.000000e+00 2.220446e-16 2.220446e-16 2.220446e-16 8.812427e-11
      8      9     10
2.220446e-16 4.488538e-10 2.220446e-16
> out=cbind(hp, wt, disp, gear, am, pred)
> out
      hp    wt  disp gear am      pred
1  110 2.620 160.0    4  1 1.000000e+00
2  110 2.875 160.0    4  1 1.000000e+00
3   93 2.320 108.0    4  1 1.000000e+00
4  110 3.215 258.0    3  0 2.220446e-16
5  175 3.440 360.0    3  0 2.220446e-16
6  105 3.460 225.0    3  0 2.220446e-16
7  245 3.570 360.0    3  0 8.812427e-11
8   62 3.190 146.7    4  0 2.220446e-16
9   95 3.150 140.8    4  0 4.488538e-10
10 123 3.440 167.6    4  0 2.220446e-16
> |
```

```
>  
> with(am.glm, null.deviance - deviance)  
[1] 43.22973  
> with(am.glm, df.null - df.residual)  
[1] 4  
> with(am.glm, pchisq(null.deviance - deviance, df.null - df.residual, lower.tail = FALSE))  
[1] 9.27207e-09  
>  
> |
```

```
>
> library(leaps)
> library(bestglm)
> X=data[,-9]
> Y=data[,9]
> Xy <- cbind(X, Y)
> bestglm(Xy, family = binomial)
Morgan-Tatar search since family is non-gaussian.
BIC
BICq equivalent for q in (0.00271755325461331, 0.849778894907533)
Best Model:
      Estimate Std. Error      z value Pr(>|z|)
(Intercept)  24.97793   211732.15  0.0001179695 0.9999059
wt          -148.46570    84415.17 -0.0017587562 0.9985967
gear         105.57256    68256.49  0.0015467037 0.9987659
There were 50 or more warnings (use warnings() to see the first 50)
>
>
```

```
> wt=mtcars$wt[1:10]
> gear=mtcars$gear[1:10]
> am=mtcars$am[1:10]
> pred=predict(am.glm, newdata, type="response")
> pred
```

	1	2	3	4	5	6	7
1.000000e+00	1.000000e+00	1.000000e+00	2.220446e-16	2.220446e-16	2.220446e-16	8.812427e-11	
	8	9	10				
2.220446e-16	4.488538e-10	2.220446e-16					

```
> out=cbind(hp, wt, am, pred)
> out
```

	hp	wt	am	pred
1	110	2.620	1	1.000000e+00
2	110	2.875	1	1.000000e+00
3	93	2.320	1	1.000000e+00
4	110	3.215	0	2.220446e-16
5	175	3.440	0	2.220446e-16
6	105	3.460	0	2.220446e-16
7	245	3.570	0	8.812427e-11
8	62	3.190	0	2.220446e-16
9	95	3.150	0	4.488538e-10
10	123	3.440	0	2.220446e-16

```
> |
```

>fitted(am.glm) ###Gives probabilities of Y=1 for all data###